


1. Compilation Phase

- **Requires:** PAPI, Extrae and BSC's LLVM
- **Input:** c/cpp + OpenMP code
 - Instrument with extrae calls and taskgraph generation pragmas
- **Output:** executable and TDG-related files



PAPI

- Performance Counters extraction API
 - Counter inside Hardware/Software Components (CPU, GPU, accelerators,...)
 - Device-dependent
- Allows us to see selected HW/SW counters around a specific code parcel
 - E.g. evaluate performance of OpenMP tasks!
- To reduce intrusion, we evaluate some perf. counters:
 - Number of CYCLES
 - L1, L2 and L3 data cache access ratios (hits vs misses)

Extræe

- Monitor and extract runtime performance traces of applications
- Designed for applications compiled and run with the shared memory model (e.g. OpenMP)
- Outputs trace files in “paraver” trace format (CSV-like)

Testing Extrae

- Let's build a small example to test extrae
- Add Extrae_init and Extrae_fini calls in the first_test project

```
cd $COURSE_PATH
cd workspace/first_test
code . #add extrae calls
make test
./test
```

Makefile

```
test: main.c
    ${CC} -I${TAFLOW_PATH}/extrae/include -L${TAFLOW_PATH}/extrae/lib main.c -fopenmp -lomptrace -o test
```

- You can also try the “heat” example provided in the exercises folder

main.c


```
#include <stdio.h>
#include <extrae.h>
int main()
{
    Extrae_init();
    #pragma omp parallel
    #pragma omp single
    {
        #pragma omp task
        printf("Hi from task 1\n");
        #pragma omp task
        printf("Hi from task 2\n");
    }
    Extrae_fini();
}
```

Testing Extrae: expected output

```
Welcome to Extrae 3.8.3    #Extrae Logging
Extrae: Detected GOMP version is 4.5
Extrae: Detected and hooked OpenMP runtime: [Intel KMPC] [GNU GOMP]
Extrae: OMP_NUM_THREADS set to 4
Extrae: Parsing the configuration file (/opt/taflow/extrae/share/example/OMP/extrae.xml)
.... #if errors with counters, just ignore them for now!
Extrae: Error! Hardware counter PAPI_TOT_INS (0x80000032) is not available. Check set 1.
... #more logging
Extrae: Tracing mode is set to: Detail.
Extrae: Successfully initiated with 1 tasks and 4 threads
      #
Hi from task 1 # application output
Hi from task 2 #
Extrae: Intermediate raw trace file created : .../test/set-0/TRACE@mypc.000008749400000000000000.mpit
...
      #more Extrae Logging
mpi2prv: Congratulations! test.prv has been generated.
```

BSC's LLVM


- C/CPP compilation with extra features in OpenMP
- Features designed for task-based OpenMP applications with dependencies
- Generates Task-Dependency Graphs (TDGs) from the OpenMP parallel regions
 - Simplifies the parallel DAG problem
 - All tasks are “siblings”
 - Connections represent dependencies
- Allows static task-to-thread mapping
- TDG is generated in
 - .dot format: ideal for analysis
 - .cpp format: ideal to apply static mapping



Example of a generated TDG

LLVM Compilation

- We can use the previous example to build the same application
- But first, let's make some changes



main.c

```
...

int main()
{
    Extrae_init();
    #pragma omp parallel
    #pragma omp single
    #pragma omp taskgraph tdg_type(static)
    {
        ...
    }
    Extrae_fini();
}
```

Makefile

```
TDG = -fopenmp-taskgraph -static-tdg -L${OMP_PATH}

test: main.c
    ${CC} ... main.c ${TDG} -fopenmp -lomptrace -o test
```


- Hints the compiler to extract a TDG

LLVM Compilation

- The compilation now generates extra files:
 - *_tdg*.dot: TDG structure as a Digraph
 - *_tdg.cpp: TDG as a C-style structure
 - tdg.hpp: contains TDG struct
- Check if these files were generated

```
...  
  
int main()  
{  
    Extrae_init();  
    #pragma omp parallel  
    #pragma omp single  
    #pragma omp taskgraph tdg_type(static)  
    {  
        ...  
    }  
    Extrae_fini();  
}
```

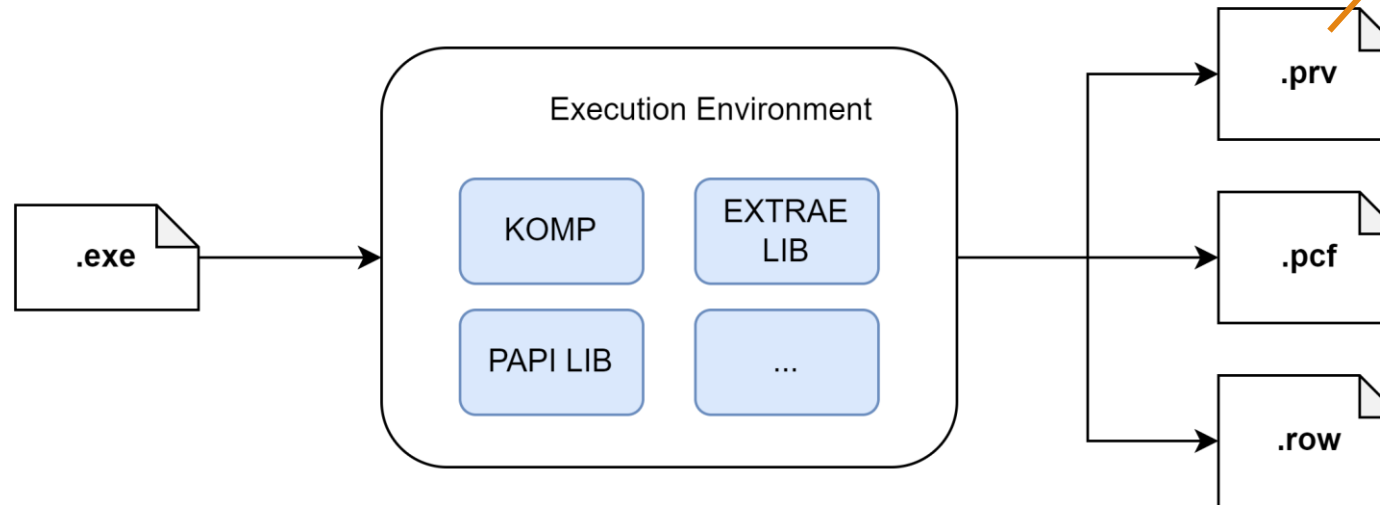
LLVM



```
1  #include "tdg.hpp"  
2  int heat_kmp_tdg_outs_1[112] = {1  
3  struct kmp_node_info heat_kmp_tdg  
4  { .static_id = 1, .task = NULL, .  
5  { .static_id = 2, .task = NULL, .  
6  { .static_id = 3, .task = NULL, .  
7  { .static_id = 4, .task = NULL, .  
8  { .static_id = 5, .task = NULL, .  
9  { .static_id = 6, .task = NULL, .  
10 { .static_id = 7, .task = NULL, .  
11 { .static_id = 8, .task = NULL, .  
12 { .static_id = 9, .task = NULL, .  
13 { .static_id = 10, .task = NULL, .
```


Running the Application

- Execution (and output) with the new LLVM is still the same

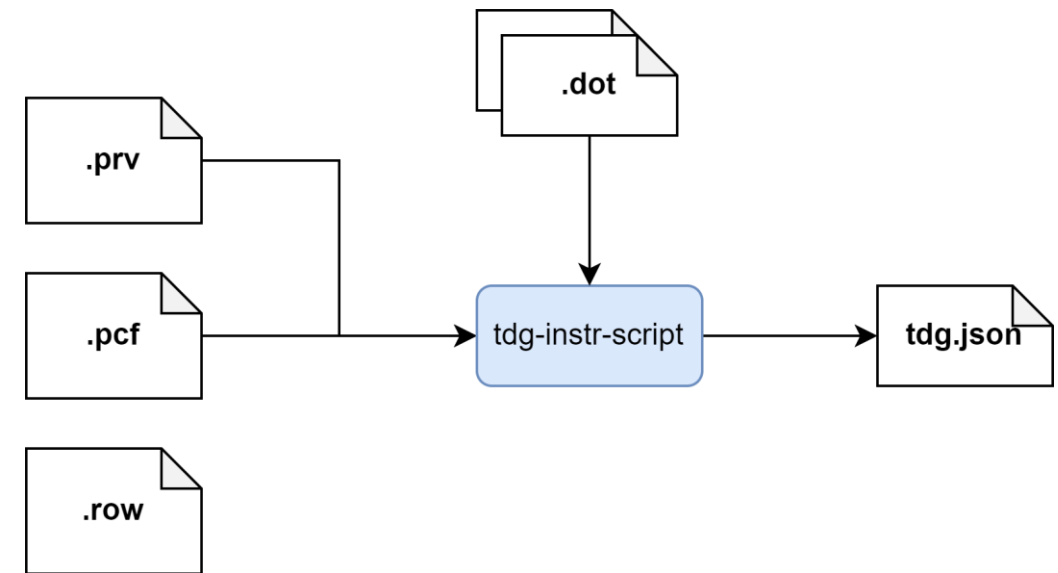


- Open this file and analyse it

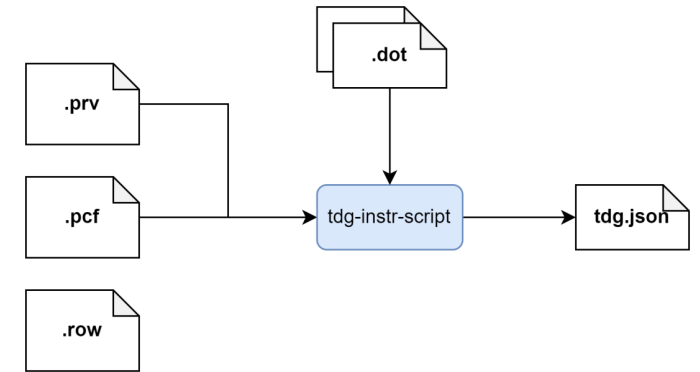
- The output trace file (prv)
 - Mainly designed for GUI-based BSC Tools (paraver, which roughly translates to “to see”)
 - Hard to read/understand (trace e.g. 2:1:1:1:1:0:40000018:1:41999999:1:42000050:0...)
 - Difficult to process/automate analysis
 - One file per execution (rewrites file)

Converting Extrae Output

- Using the python script: **tdg-instrumentation-script**
- Converts prv file into json, which contains:
 - parallel regions containing tasks (“nodes”)
 - nodes containing dependencies (“ins”, “outs”) and results
- More readable and easier to parse/manipulate



Converting Extrae Output



- Conversion (first execution): generates tdg_output.json file

```
python3 ${TAFLOW_SRC_PATH}/tdg-instrumentation-script/parsePrvAndTdg.py *.dot *.prv *.pcf
```

- Update JSON file (following executions)

```
python3 ${TAFLOW_SRC_PATH}/tdg-instrumentation-script/parsePrvAndTdg.py tdg.json *.prv *.pcf
```

TDG.json Example

```
{  
  "test": [ #application name  
    {  
      "taskgraph_id": 2658744759, #uid for the TDG  
      "nodes": {  
        "0": { #one "node" for each OpenMP Task  
          "ins": [], "outs": [], #in and out dependencies of the Task  
          "results": [  
            { #one result per execution of the OpenMP Task  
              "thread": 1,  
              "execution_begin_time": 18996782,  
              "execution_end_time": 19108462,  
              "execution_total_time": 111680,  
              "42000050": 5196, #a performance counter (see .pcf for the name)  
            }  
          ]  
        }  
      }  
    ]  
  }  
}
```

2. Profiling Phase

- A series of executions of the application
- Each execution providing a .prv file with 1 or more results
 - One for each execution of the parallel region
 - Loop in the application can be used for multiple results in same execution
- Prv file must be converted into the JSON file
 - Or results appended to an existing one
- Final result will be a JSON file
 - organized by tasks
 - each task with an array of results

2. Profiling Phase Example

- Run the following commands
 - Observe the evolution of the `tdg_output.json` file between each command

```
cd $COURSE_PATH
cd workspace/first_test
make test
./test
python3 ${TAFLOW_SRC_PATH}/tdg-instrumentation-script/parsePrvAndTdg.py *.dot *.prv *.pcf
./test
python3 ${TAFLOW_SRC_PATH}/tdg-instrumentation-script/parsePrvAndTdg.py tdg_output.json *.prv *.pcf
./test
python3 ${TAFLOW_SRC_PATH}/tdg-instrumentation-script/parsePrvAndTdg.py tdg_output.json *.prv *.pcf
...
```

- Ideally, it is better to automate this process (e.g. via a script)